

Operating Systems

Week 2 Recitation: The system call

Paul Krzyzanowski

Rutgers University

Spring 2015

System calls

- System calls are an operating system's API
 - The set of functions that the operating system exposes to processes
- If you want to the OS to do something, you tell it via a system call
- Examples

Windows	Linux
<i>NtOpenFile</i>	<i>open</i>
<i>NtReadFile</i>	<i>read</i>
<i>NtCreateProcess</i>	<i>fork</i>
<i>NtGetCurrentProcessorNumber</i>	<i>getpid</i>

See <http://j00ru.vexillium.org/ntapi/> for a list of Windows system calls
See <http://linux-documentation.com/en/man/man2/> for a list of Linux system calls

What are system calls used for?

- Anything to do with:
 - Accessing devices
 - Accessing files
 - Requesting memory
 - Setting/changing access permissions
 - Communicating with other processes
 - Stopping/starting processes
 - Setting a timer
- You need a system call to:
 - Open a file
 - Get data from the network
 - Kill a process
- You **do not** need a system call to:
 - Replace data in a string
 - Create an object (instance of a class)
 - Call a function

System calls are made via traps

- System calls request operating system services
- Operating system code executes with the processor running in *kernel* (also known as *supervisor* or *privileged*) mode
 - Privileged mode gives the CPU the rights to:
 - Execute special instructions (change interrupt masks, set hardware timers, halt the processor)
 - Access specific registers (e.g., private stack pointer)
 - Change the memory map
 - Access regions of memory that have been restricted for kernel access only
 - Access the processor's I/O ports (if the architecture has them)
- A trap takes has one parameter: index into an **Interrupt Vector Table**
 - The table is in memory that only the kernel can access
 - All addresses in the table go to well-defined entry points in the OS

Making system calls programmer-friendly

- System calls are made to look like function calls
- As a programmer, you do not want to
 - copy parameters into some special place
 - know the system call number
 - invoke a software interrupt
 - figure out how to copy any return data back
- **System call library**
 - A user-level library that is linked with your program
 - Provides a functional interface to system calls
 - Handles the work of passing parameters and getting results

System calls

Entry

Trap to system call handler

- Save state
- Verify parameters are in a valid address
- Copy them to kernel address space
- Call the function that implements the system call
 - If the function cannot be satisfied immediately then
 - Put process on a **blocked** list
 - **Context switch** to let another *ready* process run

Return from system call or interrupt

- Check for signals to the process
 - Call the appropriate handler if signal is not ignored
- Check if another process should run
 - Context switch to let the other process run
 - Put our process on a **ready** list
- Calculate time spent in the call for profiling/accounting
- Restore user process state
- Return from interrupt

The End