

Operating Systems

06. Synchronization

Paul Krzyzanowski

Rutgers University

Spring 2015

Concurrency

Concurrent threads/processes (informal)

- Two processes are concurrent if they run at the same time or if their execution is interleaved *in any order*

Asynchronous

- The processes require occasional synchronization

Independent

- They do not have any reliance on each other

Synchronous

- Frequent synchronization with each other – order of execution is guaranteed

Parallel

- Processes run at the same time on separate processors

Race Conditions

A **race condition** is a bug:

- The outcome of concurrent threads are unexpectedly dependent on a specific sequence of events.

Example

- Your current bank balance is \$1,000.
- Withdraw \$500 from an ATM machine while a \$5,000 direct deposit is coming in

Execute concurrently

Withdrawal

- Read account balance
- Subtract 500
- Write account balance

Deposit

- Read account balance
- Add 5000
- Write account balance

Possible outcomes:

Total balance = \$5500 \$500 \$6000

Synchronization


Synchronization deals with developing techniques to avoid race conditions

Something as simple as

$x = x + 1;$

Compiles to this and may cause a race condition:

```
movl  _x (%rip), %eax
addl  $1, %eax
movl  %eax, _x (%rip)
```



Potential points of preemption for a race condition

Mutual Exclusion

Critical section:

Region in a program where race conditions can arise

Mutual exclusion:

Allow only one thread to access a critical section at a time

Deadlock:

A thread is perpetually blocked (circular dependency on resources)

Starvation:

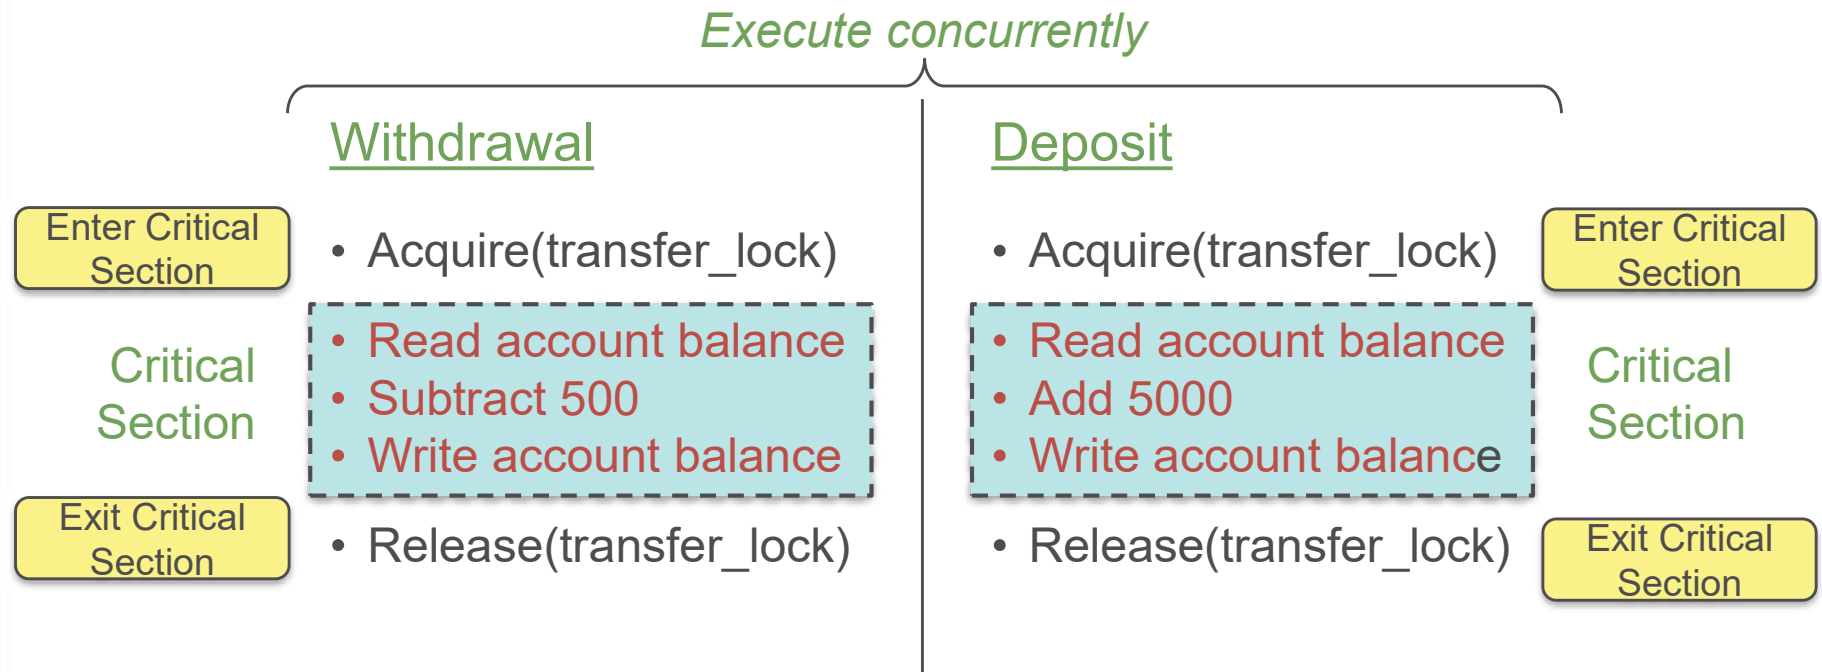
A thread is perpetually denied resources

Livelock:

Threads run but with no progress in execution

Avoid race conditions with locks

- Grab and release locks around **critical sections**
- Wait if you cannot get a lock



The Critical Section Problem

Design a protocol to allow threads to enter a critical section

Conditions for a solution

- **Mutual exclusion**: No threads may be inside the same critical sections simultaneously
- **Progress**: If no thread is executing in its critical section but one or more threads want to enter, the selection of a thread cannot be delayed indefinitely.
 - If one thread wants to enter, it should be permitted to enter.
 - If multiple threads want to enter, exactly one should be selected.
- **Bounded waiting**: No thread should wait forever to enter a critical section
- No thread running outside its critical section may block others
- A good solution will make no assumptions on:
 - No assumptions on # processors
 - No assumption on # threads/processes
 - Relative speed of each thread

Critical sections & the kernel

- Multiprocessors
 - Multiple processes on different processors may access the kernel simultaneously
 - Interrupts may occur on multiple processors simultaneously
- Preemptive kernels
 - **Preemptive kernel**: process can be preempted while running in kernel mode (the scheduler may preempt a process even if it is running in the kernel)
 - **Nonpreemptive kernel**: processes running in kernel mode cannot be preempted (but interrupts can still occur!)
- Single processor, nonpreemptive kernel
 - Free from race conditions!

Solution #1: Disable Interrupts

Disable all system interrupts before entering a critical section and re-enable them when leaving

Bad!

- Gives the thread too much control over the system
- Stops time updates and scheduling
- What if the logic in the critical section goes wrong?
- What if the critical section has a dependency on some other interrupt, thread, or system call?
- What about multiple processors? Disabling interrupts affects just one processor

Advantage

- Simple, guaranteed to work
- Was often used in the uniprocessor kernels

Solution #2: Software Test & Set Locks

Keep a shared lock variable:

```
while (locked) ;  
locked = 1;  
/* do critical section */  
locked = 0;
```

Disadvantage:

- Buggy! There's a race condition in setting the lock

Advantage:

- Simple to understand. It's been used for things such as locking mailbox files

Solution #3: Lockstep Synchronization

Take turns

Thread 0

```
while (turn != 0);  
critical_section();  
turn = 1;
```

Thread 1

```
while (turn != 1);  
critical_section();  
turn = 0;
```

Disadvantage:

- Forces strict alternation; if thread 2 is really slow, thread 1 is slowed down with it. *Turns asynchronous threads into synchronous threads*

Software solutions for mutual exclusion

- Peterson's solution (page 207 of text) , Dekker's, & others
- Disadvantages:
 - Difficult to implement correctly
Have to rely on `volatile` data types to ensure that compilers don't make the wrong optimizations
 - Difficult to implement for an arbitrary number of threads

Let's turn to hardware for help

Help from the processor

Atomic (indivisible) CPU instructions that help us get locks

- Test-and-set
- Compare-and-swap
- Fetch-and-Increment

These instructions execute in their entirety: they cannot be interrupted or preempted partway through their execution

Test & Set

Set the lock but get told if it already was set (in which case you don't have it)

```
ATOMIC {  
    int test_and_set(int *x) {  
        last_value = *x;  
        *x = 1;  
        return last_value;  
    }  
}
```

How you use it to lock a critical section (i.e., enforce mutual exclusion):

```
while (test_and_set(&lock) == 1) ;    /* spin */  
/* do critical section */  
lock = 0;    /* release the lock */
```


Fetch & Increment

Increment a memory location; return previous value

ATOMIC {

```
int fetch_and_increment(int *x) {  
    last_value = *x;  
    *x = *x + 1;  
    return last_value;  
}
```

The problem with spin locks

- All these solutions require busy waiting
 - Tight loop that spins waiting for a turn: busy waiting or spin lock
- Nothing useful gets done!
 - Wastes CPU cycles

Priority Inversion

- Spin locks may lead to **priority inversion**
- The process with the lock may not be allowed to run!
 - Suppose a lower priority process obtained a lock
 - Higher priority process is always ready to run but loops on trying to get the lock
 - Scheduler always schedules the higher-priority process
 - **Priority inversion**
 - If the low priority process would get to run & release its lock, it would then accelerate the time for the high priority process to get a chance to get the lock and do useful work
 - Try explaining that to a scheduler!

Spin locks aren't great

Can we block until we can get the critical section?

Semaphores

- Count # of wake-ups saved for future use
- Two atomic operations:

```
down(sem s) {  
    if (s > 0)  
        s = s - 1;  
    else  
        sleep on event s  
}
```

```
up(sem s) {  
    if (someone is waiting on s)  
        wake up one of the threads  
    else  
        s = s + 1;  
}
```

```
//initialize  
mutex = 1;  
  
down(&mutex)  
  
// critical section  
  
up(&mutex)
```

Binary semaphore

Semaphores

Count the number of threads that may enter a critical section at any given time.

- Each *down* decreases the number of future accesses
- When no more are allowed, processes have to wait
- Each *up* lets a waiting process get in

Producer-Consumer example

- Producer
 - Generates items that go into a buffer
 - Maximum buffer capacity = N
 - If the producer fills the buffer, it must wait (sleep)
- Consumer
 - Consumes things from the buffer
 - If there's nothing in the buffer, it must wait (sleep)
- This is known as the *Bounded-Buffer Problem*

Producer-Consumer example

```
sem mutex=1, empty=N, full=0;
producer() {
    for (;;) {
        produce_item(&item);    // produce something
        down(&empty);           // decrement empty count
        down(&mutex);            // start critical section
        enter_item(item);        // put item in buffer
        up(&mutex);              // end critical section
        up(&full);               // +1 full slot
    }
}
consumer() {
    for (;;) {
        down(&full);            // one less item
        down(&mutex);           // start critical section
        remove_item(item);      // get the item from the buffer
        up(&mutex);              // end critical section
        up(&empty);             // one more empty slot
        consume_item(item);     // consume it
    }
}
```


Condition Variables / Monitors

- Higher-level synchronization primitive
- Implemented by the programming language / APIs
- Two operations:
 - wait(*condition_variable*)
 - Block until *condition_variable* is “signaled”
 - signal(*condition_variable*)
 - Wake up one process that is waiting on the condition variable
 - Also called notify

Synchronization

Part II: Inter-Process Message Passing

Communicating processes

- Must:
 - Synchronize
 - Exchange data
- Message passing offers:
 - Data communication
 - Synchronization (via waiting for messages)
 - Works with processes on different machines

Message passing


- Two primitives:
 - send(destination, message)
 - receive(source, message)
- Operations may or may not be blocking

Messaging: Rendezvous

- Sending process blocked until receive occurs
- Receive blocks until a send occurs
- Advantages:
 - No need for message buffering if on same system
 - Easy & efficient to implement
 - Allows for tight synchronization
- Disadvantage:
 - Forces sender & receiver to run in lockstep

Messaging: Direct Addressing

- Sending process identifies receiving process
- Receiving process can identify sending process
 - Or can receive it as a parameter




Messaging: Indirect Addressing

- Messages sent to an intermediary data structure of FIFO queues
- Each queue is a mailbox
- Simplifies multiple readers


Mailboxes




Single sender, single reader



Single sender, multiple readers



Multiple senders, single reader



Multiple senders, multiple readers

Other common IPC mechanisms

- Shared files
 - File locking allows concurrent access control
 - Mandatory or advisory
- Signal
 - A simple poke
- Pipe
 - Two-way data stream using file descriptors (but not names)
 - Need a common parent or threads in the same process
- Named pipe (FIFO file)
 - Like a pipe but opened like a file
- Shared memory

Conditions for deadlock

Four conditions must hold

1. **Mutual exclusion**
 - Only one thread can access a critical section (resource) at a time
2. **Hold and wait**
 - A thread holds a resource but waits for another resource
3. **Non-preemption of resources**
 - Resources can only be released voluntarily
4. **Circular wait**
 - There is a cyclic dependency of threads waiting on resources

Deadlock

- Resource allocation

- Resource R_1 is allocated to process P_1 : *assignment edge*




- Resource R_1 is requested by process P_1 : *request edge*



- **Deadlock** is present when the graph has **cycles**

Deadlock example



Circular dependency among four processes and four resources leads to deadlock

Dealing with deadlock

- **Deadlock prevention**
 - Ensure that at least one of the necessary conditions cannot hold
- **Deadlock avoidance**
 - Provide advance information to the OS on which resources a process will request.
 - OS can then decide if the process should wait
 - *But knowing which resources will be used (and when) is hard!*
(impossible, really)
- **Deadlock detection**
 - Detect when a deadlock occurs and then deal with it
- **Ignore the problem**
 - Let the user deal with it (most common approach)

The End